

1. Provider: OpenAI

Model: GPT-5

Release Date: August 7, 2025

Knowledge Cutoff: August 31, 2025

Context Window: 400,000 tokens

Cost of Tokens: \$2.50 / 1M input · \$15.00 / 1M output

Model Type:  Private / Commercial

2. Provider: Anthropic

Model: Claude Opus 4.6

Release Date: February 5, 2026

Knowledge Cutoff: Early 2026 (exact date not publicly disclosed)

Context Window: 1,000,000 tokens (1M)

Cost of Tokens: Premium tier (exact rate follows Claude API pricing structure)

Model Type:  Private / Commercial

3. Provider: Google DeepMind

Model: Gemini 2.5 Pro

Release Date: Early 2025

Knowledge Cutoff: Early 2025

Context Window: 1,000,000 – 2,000,000 tokens

Cost of Tokens: ~\$2.00 / 1M input · \$12.00 / 1M output

Model Type:  Private / Commercial

4. Provider: Meta

Model: Llama 4 Scout

Release Date: April 5, 2025

Knowledge Cutoff: August 2024

Context Window: 10,000,000 tokens (10M — largest available)

Cost of Tokens: Free (open weights); API ~\$0.19 / 1M tokens (blended)

Model Type:  Public / Open Source (Custom Meta License)

5. Provider: DeepSeek

Model: DeepSeek V3

Release Date: January 2025

Knowledge Cutoff: Late 2024

Context Window: 128,000 tokens

Cost of Tokens: \$0.55 / 1M input · \$2.19 / 1M output

Model Type:  Public / Open Source (MIT License)

6. Provider: Mistral AI

Model: Mistral Large 3

Release Date: December 2025

Knowledge Cutoff: Late 2025

Context Window: 256,000 tokens

Cost of Tokens: \$0.50 / 1M input · \$1.50 / 1M output

Model Type:  Public / Open Source (Apache 2.0)

7. Provider: xAI

Model: Grok 4

Release Date: September 19, 2025

Knowledge Cutoff: December 2025 (+ live web search capability)

Context Window: 1,000,000 tokens

Cost of Tokens: \$3.00 / 1M input · \$15.00 / 1M output

Model Type:  Private / Commercial

8. Provider: Cohere

Model: Command A

Release Date: March 2026

Knowledge Cutoff: Early 2026

Context Window: 256,000 tokens

Cost of Tokens: \$2.50 / 1M input · \$10.00 / 1M output

Model Type:  Private / Commercial

9. Provider: Alibaba (Qwen Team)

Model: Qwen 3 (Qwen3-Max)

Release Date: April 28, 2025

Knowledge Cutoff: Early 2025

Context Window: Up to 1,000,000 tokens (native 256K, extendable)

Cost of Tokens: From \$1.04 / 1M input tokens (via Alibaba Cloud)

Model Type:  Public (open weights) +  Commercial API

10. Provider: Google DeepMind

Model: Gemma 4 (31B Dense)

Release Date: April 2, 2026

Knowledge Cutoff: Early 2026

Context Window: 256,000 tokens

Cost of Tokens: Free (open weights, self-hosted)

Model Type:  Public / Open Source (Apache 2.0)